
Adaptive Loss Weighting Is Costing You an Order of Magnitude: Physics-Informed Networks and Classical Solvers on the Heterogeneous Elastic Wave Equation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Causal time weighting and GradNorm loss balancing are standard equipment in
2 modern physics-informed neural network (PINN) training. On the 1D elastic wave
3 equation in heterogeneous media we find that, run exactly as specified, **both hurt**:
4 switching them off beats the full recipe in every one of the 21 architecture–material
5 cells we measured, by $1.8\text{--}220\times$ (median $17\times$), and takes the best network from
6 2.30% to 0.302%. Under the recipe, 29 of 42 runs score at or above the 100% that
7 predicting zero earns; without it, none do. The mechanism is measurable rather
8 than mysterious: in 34 of 36 causal runs the schedule’s frontier stalls, leaving
9 part of the time domain weighted at exactly zero for the whole of training, and
10 GradNorm independently settles at a weighting that puts the PDE residual three
11 orders of magnitude below the boundary term. Only 2 of those runs open all sixteen
12 slabs: causal weighting is conditional on the optimiser reaching the early slabs, and
13 when that fails it is much worse than not weighting at all. We reach this by way of
14 a control the PINN literature rarely runs: a characterised classical baseline on the
15 same problem, same grid, same metric. Building one turns up a second result —
16 second-order finite differences and linear finite elements both stall at 0.039% on
17 all three materials, a floor set by the first-order absorbing boundary rather than the
18 interior stencil, which any PINN scored against such a reference silently inherits.
19 Even after the ablation, the best network’s error is $7.5\times$ that of a finite-difference
20 solve costing 0.0177 s on one CPU core, so the headline is not that PINNs have
21 been fixed. It is that a widely used training recipe is being credited for accuracy
22 it is in fact removing, and that the diagnostic which reveals this ($w_{\min}$, already
23 computed by every causal implementation) costs nothing to report.

24 1 Introduction

25 Physics-informed neural networks [Raissi et al., 2019] solve a PDE by minimising its residual at
26 collocation points, and a large literature now supplies training machinery to make that minimisation
27 work: causal time weighting [Wang et al., 2022], gradient-norm balancing of the loss terms [Wang
28 et al., 2021a, 2023], residual-based attention [Anagnostopoulos et al., 2023], Fourier feature embed-
29 dings [Tancik et al., 2020, Wang et al., 2021b], residual-adaptive architectures [Wang et al., 2024],
30 Kolmogorov–Arnold layers [Liu et al., 2024, Bozorgasl and Chen, 2024, Shukla et al., 2024], and
31 second-order optimisers [Vyas et al., 2024, Wang et al., 2025]. Each is reported to improve on a
32 baseline PINN. What is much less often reported is how the resulting system compares to a classical
33 discretisation of the same equation, run on the same problem, and measured with the same metric.

34 This paper makes that comparison on one deliberately modest problem: the 1D elastic wave equation
 35 in a heterogeneous medium,

$$\rho(x) u_{tt} = \partial_x [E(x) u_x], \quad x \in [x_{\min}, x_{\max}], \quad t \in [0, 1], \quad (1)$$

36 with a Gaussian-derivative initial pulse, zero initial velocity, and first-order radiation conditions at
 37 both ends. The problem is small enough that a classical solver is essentially free, which is exactly
 38 what makes it a useful control: any claim that a PINN is doing something valuable has to be made
 39 against a baseline whose accuracy and cost are both known.

40 Running that control produced a result we were not looking for. Two of the adaptive schemes above —
 41 causal time weighting and GradNorm balancing — turn out to *cost* an order of magnitude of accuracy
 42 on this problem when run as specified. Switching both off improves relative L^2 by $1.8\text{--}220\times$. Neither
 43 is misbehaving: both do exactly what their specification says, and what the specification does here
 44 is leave most of the time domain untrained and the PDE residual weighted a thousand times below
 45 the boundary term. This is invisible if a run is reported by its final error alone, and obvious from
 46 diagnostics the methods already compute.

47 Contributions.

- 48 1. **A characterised classical baseline.** We implement three discretisations of (1) — second-
 49 order finite differences in flux form, linear finite elements with a lumped mass matrix,
 50 and Chebyshev collocation on the velocity–stress system — and measure error against a
 51 self-converged spectral solution. The spectral solver reaches $3.6e\text{--}5\%$ relative L^2 on the
 52 homogeneous case in 0.1303s; the two second-order schemes are indistinguishable from
 53 each other and *plateau* at 0.039%.
- 54 2. **An accuracy floor in the reference itself.** That plateau is not the interior stencil. It is the
 55 first-order Mur absorbing boundary [Mur, 1981]: refining from 1 600 to 6 400 cells moves
 56 the error by less than 0.1% of its own value. Every PINN number in this literature that is
 57 scored against a finite-difference reference inherits this floor, and it is worth knowing where
 58 it sits before reading a small error as a small error.
- 59 3. **The published training recipe, run end to end.** We evaluate seven coordinate-network
 60 architectures on three material profiles under two initial condition treatments, with the
 61 causal/GradNorm recipe as specified, plus SOAP and residual-based-attention arms.
- 62 4. **Both adaptive weighting schemes make things worse.** Switching causal weighting and
 63 GradNorm off — pinning $w_{\text{pde}} = w_{\text{bc}} = w_{\text{ic}} = 1$ and weighting every time slab equally —
 64 improves relative L^2 in 21 of 21 architecture–material cells, which is all of them, by 1.8--
 65 $220\times$ (median $17\times$), from 2.30% to 0.302% for the best architecture. This is the paper’s
 66 main result and it was not what we set out to look for.
- 67 5. **A mechanism, not a mystery.** The causal frontier advances two slabs of sixteen and
 68 stops, so most of the time domain carries weight exactly zero for all 24 000 Adam steps
 69 including the continuation budget; GradNorm independently settles at $w_{\text{pde}} \approx 0.001$ against
 70 $w_{\text{bc}} \approx 1.999$. Both are visible in diagnostics these methods already compute.

71 2 Problem, materials, and what “exact” means here

72 **Equation and non-dimensionalisation.** We use the conservative form (1) throughout. Expanding it
 73 gives $\rho u_{tt} = E u_{xx} + E_x u_x$, so the frequently used $u_{tt} = c(x)^2 u_{xx}$ drops the $E_x u_x$ term — precisely
 74 the term that sets reflection and transmission amplitudes at a material interface. Each material is
 75 non-dimensionalised by a reference stiffness E_{ref} and density ρ_{ref} and a half-length L_{ref} , so that
 76 $\rho \equiv 1$, $E = O(1)$, and $x \in [-1, 1]$.

77 **Materials.** Three profiles of increasing difficulty, all with $\rho \equiv 1$ so that $c(x) = \sqrt{E(x)}$: *Homoge-*
 78 *neous* ($E \equiv 1$); *TwoLayer*, a tanh-smoothed step from $E = 1$ to $E = 1.5$ at $x = 0$ with interface
 79 half-width 0.02; and *MultiLayer*, six tanh-blended layers with E from 1 to 2.5. The interfaces are
 80 smooth but narrow: at half-width 0.02 on a domain of length 2, resolving one costs roughly a hundred
 81 grid points.

Table 1: Cost of accuracy for the three classical discretisations, measured against a converged spectral solution. Each pair of columns gives the cheapest resolution on the refinement ladder that reaches the target and the wall-clock seconds it took (best of three, one CPU core). *n/a* means no resolution on the ladder reaches it. **Floor** is where the method stops improving: the two second-order schemes stall at the same 0.04% on all three materials, which is the first-order Mur absorbing boundary rather than the interior stencil.

material	method	$\leq 1\%$		$\leq 0.1\%$		$\leq 0.01\%$		floor
		<i>n</i>	s	<i>n</i>	s	<i>n</i>	s	%
Homogeneous	f d2 (FD2)	200	0.0020	800	0.0062	<i>n/a</i>	–	0.040
	fem_p1 (P1 FEM)	200	0.0021	800	0.0070	<i>n/a</i>	–	0.040
	cheb (spectral)	64	0.1303	64	0.1303	64	0.1303	$\approx$ ref.
TwoLayer	f d2 (FD2)	200	0.0022	800	0.0076	<i>n/a</i>	–	0.040
	fem_p1 (P1 FEM)	200	0.0023	800	0.0081	<i>n/a</i>	–	0.040
	cheb (spectral)	96	0.3154	128	0.6469	192	2.05	9.5e-4
MultiLayer	f d2 (FD2)	200	0.0022	800	0.0101	<i>n/a</i>	–	0.039
	fem_p1 (P1 FEM)	200	0.0023	800	0.0106	<i>n/a</i>	–	0.039
	cheb (spectral)	64	0.2167	128	0.9383	256	6.84	2.7e-3

Initial and boundary conditions. $u(x, 0) = f(x)$ is the derivative of a Gaussian of width $\sigma_g = 0.1$, normalised to unit peak, and $u_t(x, 0) = 0$. Because the initial velocity vanishes, the pulse splits into two counter-propagating waves of roughly half amplitude. Both boundaries carry the first-order radiation condition $u_t \mp c u_x = 0$, so waves leave the domain rather than reflecting.

Metric. All errors are relative L^2 in percent on a shared 201×101 evaluation grid over (x, t) , $100 \|\hat{u} - u\|_2 / \|u\|_2$. This has a useful calibration property: a model that predicts zero everywhere scores approximately 100%, so any number near or above 100 indicates a solution that carries no signal.

3 Classical baselines

We implement three discretisations of (1), all sharing the problem above and all returning u resampled onto the evaluation grid so that the metric never depends on a solver’s internal resolution.

f d2 is the explicit second-order leapfrog scheme in flux form, with harmonic-mean interface stiffness $E_{i+1/2} = 2E_i E_{i+1} / (E_i + E_{i+1})$ and Mur-discretised absorbing boundaries $u_0^{n+1} = u_1^n + \beta(u_1^{n+1} - u_0^n)$, $\beta = (c \Delta t - \Delta x) / (c \Delta t + \Delta x)$. This is the reference solver the PINN runs are scored against.

fem_p1 uses linear elements with two-point Gauss evaluation of the element coefficients, a lumped mass matrix, and explicit central differences. The radiation condition appears in the weak form as a boundary ρc dashpot, so it needs no special discretisation.

cheb collocates the velocity–stress system $\rho v_t = \partial_x \sigma$, $\sigma_t = E \partial_x v$, $u_t = v$ on Chebyshev–Gauss–Lobatto nodes and integrates with RK4. Splitting into first order is what makes the boundary tractable: the radiation condition becomes “the incoming Riemann invariant is zero”, imposed as a projection at every Runge–Kutta stage. The second-order form has no stable second-order absorbing update — advancing u at the end nodes from u^{n-1} decouples them from the interior and the corner modes grow without bound, which is what we observed before switching formulations.

Ground truth. We take a converged cheb solution ($n = 512$) as ground truth rather than a fine finite-difference solve, for the reason Table 1 makes plain. Its self-convergence is reported alongside every result: on *Homogeneous*, $n = 384$ and $n = 512$ differ by 1.3×10^{-12} percent.

The two second-order schemes are the same scheme. f d2 and fem_p1 agree to within 2% of their own error at every resolution on every material. This is expected rather than surprising — a lumped-mass P1 element on a uniform grid reduces to the three-point stencil — but it is a useful control: two independent implementations, one derived from a Taylor expansion and one from a

Table 2: The finite-difference reference the PINN runs are scored against, measured against converged spectral ground truth. Refinement stops paying at roughly $N_x = 1600$; beyond that the error is set by the first-order absorbing boundary. $N_x = 1000$ is the setting the sweep uses.

material	relative L^2 (%) of fd2 at N_x					
	250	500	1000	2000	4000	8000
Homogeneous	0.351	0.071	0.040	0.040	0.040	0.040
TwoLayer	0.502	0.138	0.046	0.040	0.040	0.040
MultiLayer	0.622	0.146	0.048	0.040	0.039	0.040

weak form with a different boundary treatment, landing on the same numbers means neither carries an implementation error large enough to matter at this scale.

The reference has a floor, and it is the boundary. Both second-order schemes stop improving at 0.039%, and Table 2 shows where: going from $N_x = 1000$ to $N_x = 8000$ — a $64\times$ increase in work — changes the error by less than 2% of its own value. The floor is identical across the three materials to three significant figures, which is what identifies it: the interior physics differs enormously between a constant medium and a six-layer stack, so a shared limit cannot come from the interior stencil. It comes from the first-order Mur absorbing boundary, which admits a small reflected wave that no amount of grid refinement removes. The spectral solver, whose radiation condition is exact on the characteristic variables, has no such floor and self-converges to $1.3e-12\%$.

This matters beyond the classical comparison. The PINN runs in Section 5 are scored against fd2 at $N_x = 1000$, which is 0.040% from the true solution. That is far below any network error we measured, so no conclusion here is threatened by it — but it does bound what a PINN result on this problem could ever mean, and reporting a PINN error of 10^{-4} against such a reference would be reporting agreement with a boundary artefact.

Cost. Reaching 0.1% costs 6 ms to 11 ms on one CPU core for the second-order schemes, and 0.01% is unreachable for them at any resolution. The spectral solver reaches $9.5e-4\%$ on *TwoLayer* in 5.06 s and $2.7e-3\%$ on *MultiLayer* in 6.84 s. On *Homogeneous* it reaches $3.6e-5\%$ in 0.1303 s at $n = 64$. These are the numbers a network has to be compared against.

4 Coordinate networks

Architectures. Seven, spanning the two families the PINN literature has converged on. *VanillaPINN* is a tanh MLP. *FourierFeaturePINN* prepends a trainable random Fourier embedding [Tancik et al., 2020, Wang et al., 2021b]. *PirateNet* uses a fixed Fourier embedding, random weight factorisation, and residual blocks whose gate is initialised to zero [Wang et al., 2024]. On the Kolmogorov–Arnold side [Liu et al., 2024]: *KAN* places a B-spline on each edge together with a SiLU base branch; *WavKAN* replaces the spline with a wavelet [Bozorgasl and Chen, 2024]; *ChebyshevKAN* is the cPIKAN basis, a tanh-squashed Chebyshev recurrence per edge [Shukla et al., 2024]; and *FourierWavKAN* feeds a random Fourier embedding into a WavKAN trunk.

Loss. Writing $r = \rho u_{tt} - \partial_x[Eu_x]$ for the residual, the objective is

$$\mathcal{L} = w_{\text{pde}} \mathcal{L}_{\text{pde}} + w_{\text{bc}} \mathcal{L}_{\text{bc}} + w_{\text{ic}} \mathcal{L}_{\text{ic}}, \quad (2)$$

on a 100×100 collocation grid. Two adaptive schemes set the pieces.

Causal weighting [Wang et al., 2022] splits $[0, 1]$ into $M = 16$ time slabs and weights slab m by $\omega_m = \exp(-\varepsilon \sum_{k < m} \mathcal{L}_k)$, so a slab only receives gradient once the slabs before it are converged. We report $w_{\text{min}} = \min_m \omega_m$, the standard progress diagnostic: it reaches 1 when the whole time domain is being trained. The tolerance ε is annealed from 0.1 to 1.0, and a continuation loop grants up to 9000 extra Adam steps while $w_{\text{min}} < 0.9$.

GradNorm balancing [Wang et al., 2021a, 2023] rescales $w_{\text{pde}}, w_{\text{bc}}, w_{\text{ic}}$ every 100 steps so that each term contributes a comparable gradient norm, with exponential smoothing $\alpha = 0.9$.

Table 3: Coordinate networks under the published recipe (causal weighting + GradNorm, Adam then L-BFGS). Relative L^2 (%) against the finite-difference reference, taking the better reported Adam or L-BFGS checkpoint; **hard** is the constraint ansatz, **soft** the IC loss. A zero field scores ≈ 100 . For scale, the classical column of Table 1 reaches 0.04% in under 20 ms.

architecture	Homogeneous		TwoLayer		MultiLayer	
	hard	soft	hard	soft	hard	soft
VanillaPINN	129	178	153	182	149	189
FourierFeature	2.30	181	8.93	177	11.9	174
PirateNet	105	183	104	193	105	203
KAN (spline)	123	125	95.9	321	97.2	364
WavKAN	24.9	145	60.2	172	63.9	165
ChebyshevKAN	37.9	138	13.2	135	23.0	131
FourierWavKAN	120	130	18.2	138	13.9	165

Initial condition, two ways. Either a hard constraint, $\hat{u} = f(x)e^{-(15t)^2/2} + \tanh^2(25t)\mathcal{N}(x, t)$, which satisfies the IC exactly and removes $\mathcal{L}_{ic}$ from (2); or a soft IC loss with a warm-up phase and a fixed pre-scale that keeps the IC gradient from being starved early on.

Optimisation. Two phases: 15 000 Adam steps [Kingma and Ba, 2015] at 10^{-3} , then 2 000 L-BFGS steps [Liu and Nocedal, 1989]. L-BFGS runs with causal weighting disabled, because weights that move inside the closure corrupt the line-search objective. Model selection uses a stationary validation metric (unweighted per-slab residual plus BC and IC terms), not the weighted training loss, which moves as the weights move. Variant arms replace Adam with SOAP [Vyas et al., 2024, Wang et al., 2025] or causal weighting with residual-based attention [Anagnostopoulos et al., 2023].

Implementation. JAX [Bradbury et al., 2018] with Equinox [Kidger and Garcia, 2021]; second derivatives are taken as nested `jax.grad` over a single-sample model and vectorised with `jax.vmap`, which avoids reverse-over-reverse accumulation. Runs execute on NVIDIA H100 NVL GPUs.

5 Results

How the initial condition is imposed dominates everything else. Across the architectures in Table 3, moving the initial condition from a soft loss term to the hard-constraint ansatz is worth more than any choice of architecture. Soft-IC runs sit at or above 100% — the score of the zero field — meaning the trained network carries essentially no signal, while the same architecture with the ansatz reaches single or low double digits. The IC loss competes with the residual for gradient budget throughout training, and on this problem it loses; the ansatz removes the competition by construction.

Even the best network remains behind a classical solve. Under the published recipe the strongest configuration reaches 2.30%; the best anywhere in the sweep, an RBA run, reaches 1.34%; and the best after the ablation of Section 5 reaches 0.302%. The finite-difference reference reaches 0.040% in 0.0177 s on one CPU core, and the spectral solver goes further still. Figure 1 places both families on the same axes. Nothing about this is specific to a weak baseline: the classical methods here are textbook schemes, not tuned ones.

The causal frontier almost never leaves the first slabs. Of the 36 runs that use a causal schedule, 34 end training with at least one slab still weighted at exactly zero, and 5 of them never open a second slab at all. Their median error is 96.0%. The 9 000-step continuation the runner grants while $w_{\min} < 0.9$ does not rescue them: it is spent, and the frontier does not move.

Only 2 of the 36 runs opened all sixteen. Figure 2 puts a strong stalled run next to the run that advances furthest, so the stall cannot be dismissed as a generally bad configuration. The latter opens 16 of sixteen slabs. This diagnostic is an association, not evidence of a mechanism — the evidence is the ablation below. It nevertheless sharpens the claim: causal weighting is conditional on the optimiser driving the early slabs to tolerance, and on this problem that almost never happens. When it fails to happen, the schedule is far worse than not having used one.




Figure 1: What accuracy costs. Wall-clock seconds against relative L^2 , both logarithmic. Connected markers are the classical refinement ladders (CPU, single core); hollow markers are the spectral solver agreeing with the reference to its own precision rather than a measured error. Diamonds are individual PINN runs at their measured training time on an H100. The dotted line at 100% is the score of predicting zero. The two families are separated by roughly four orders of magnitude in time and two to three in error.




Figure 2: Causal weight ω_m of each of the 16 time slabs across the Adam phase, with each panel’s final relative L^2 in its title. *Left*: the strongest distinct run whose frontier stalls — the schedule opens a handful of slabs and stops, and the rest are weighted at exactly zero for the whole of training, continuation steps included. *Right*: the run that advances furthest, opening 16 of sixteen slabs.

184 **Both adaptive schemes are costing accuracy.** Table 4 is the main result. Removing *both* schemes
185 — pinning $w_{\text{pde}} = w_{\text{bc}} = w_{\text{ic}} = 1$ and collapsing the causal schedule to a single slab — wins in
186 21 of the 21 architecture–material cells measured against the published recipe, by $1.8\text{--}220\times$ with a
187 median of $17\times$. The Fourier-feature network goes from 2.30% under the recipe to 0.302% without it;
188 the plain MLP on *TwoLayer* likewise goes from worse than predicting zero to a single-digit error.
189 Nothing was added to obtain this: same architecture, same step budget, same optimiser, two adaptive
190 schemes switched off.

191 The change in the shape of the distribution is as striking as the change in its best value. Under the
192 recipe, 29 of 42 runs score at or above 100%, meaning the trained network carries no more signal

Table 4: Switching off the two adaptive weighting schemes, hard-constraint IC. Relative L^2 (%) at the better reported Adam or L-BFGS checkpoint; *causal+GradNorm* is the published recipe and *both off* removes both adaptive schemes.

material	architecture	causal+GradNorm	GradNorm off	causal off	both off
Homogeneous	VanillaPINN	129	12.9	7.37	1.16
	FourierFeature	2.30	79.3	1.000	0.302
	PirateNet	105	—	3.23	3.45
	KAN (spline)	123	—	437	0.560
	WavKAN	24.9	—	21.3	2.70
	ChebyshevKAN	37.9	—	6.40	1.62
	FourierWavKAN	120	—	2.31	0.624
TwoLayer	VanillaPINN	153	—	6.49	1.42
	FourierFeature	8.93	—	1.01	0.554
	PirateNet	104	—	6.01	1.14
	KAN (spline)	95.9	—	112	0.836
	WavKAN	60.2	—	28.7	5.26
	ChebyshevKAN	13.2	—	10.4	3.16
	FourierWavKAN	18.2	—	4.13	0.923
MultiLayer	VanillaPINN	149	—	19.9	16.2
	FourierFeature	11.9	—	2.17	1.51
	PirateNet	105	—	4.18	2.76
	KAN (spline)	97.2	—	207	5.71
	WavKAN	63.9	—	23.0	16.5
	ChebyshevKAN	23.0	—	10.9	12.6
	FourierWavKAN	13.9	—	7.69	5.55

than a zero field. With both schemes off, every cell we have measured lands well below that threshold. Architectures that looked hopeless under the recipe become ordinary once the weighting is removed.

The two columns in between separate the causes. Dropping causal weighting alone accounts for most of the gain in the representative MLP and Fourier-feature rows. This is consistent with Figure 2: a schedule that leaves later slabs closed is training on a fraction of the domain, and equal weighting at least trains all of it. Dropping GradNorm alone is not a uniform fix, but once causal weighting is gone, removing it helps both representative rows. Under the full recipe GradNorm can settle with the PDE residual weighted orders of magnitude below the boundary term.

The causal frontier stalls with GradNorm on *and* off, so the stall is a property of the schedule on this problem rather than a downstream consequence of the weighting. Both schemes are behaving exactly as specified; the specification is what does not suit this problem.

The RBA arm corroborates this independently. Table 5 reports the SOAP optimiser, residual-based attention (RBA) in place of causal weighting, and an L-BFGS-only schedule with no first-order phase. The entry that matters here is RBA, because it is the one variant that *removes* causal weighting: it replaces the time-slab schedule with pointwise residual weights and reports $w_{\min} = 1$ by construction. RBA runs take the top places in the entire sweep, reaching 1.34% against the recipe’s 2.30% on the same architecture.

That is worth reading carefully, because the natural attribution is wrong. One would credit the gain to the attention mechanism. The ablation says otherwise: weighting every time equally, with no attention at all, does better still (0.302%). On this problem, most of what RBA buys is not what it adds but what it displaces.

6 Discussion

We set out to run a control — a classical baseline on the same problem, same grid, same metric — and the control turned into the finding. Two adaptive schemes that are close to standard equipment in PINN training were, on this problem, removing an order of magnitude of accuracy while being credited with producing it.

Table 5: Variant training arms, hard-constraint IC, relative L^2 (%) at the better reported first- or second-order checkpoint. SOAP replaces Adam; RBA replaces causal weighting with residual-based attention; L-BFGS-only removes the first-order phase entirely.

material	architecture	base	SOAP	RBA	SOAP+RBA	lbfgs_only
Homogeneous	VanillaPINN	129	12.1	6.91	—	92.8
Homogeneous	FourierFeature	2.30	339	1.34	6.89	6.93
Homogeneous	PirateNet	105	0.353	17.8	—	8.24
Homogeneous	KAN (spline)	123	—	—	—	93.1
Homogeneous	WavKAN	24.9	238	9.54	—	96.8
Homogeneous	ChebyshevKAN	37.9	19.7	11.4	6.74	40.4
Homogeneous	FourierWavKAN	120	—	—	—	16.8
TwoLayer	VanillaPINN	153	11.3	4.74	—	93.3
TwoLayer	FourierFeature	8.93	7648	2.32	15.7	146
TwoLayer	PirateNet	104	1.46	1.99	—	6.75
TwoLayer	KAN (spline)	95.9	—	—	—	93.7
TwoLayer	WavKAN	60.2	281	23.7	—	98.5
TwoLayer	ChebyshevKAN	13.2	23.2	8.64	4.62	17.3
TwoLayer	FourierWavKAN	18.2	—	—	—	29.7
MultiLayer	VanillaPINN	149	96.0	15.7	—	94.7
MultiLayer	FourierFeature	11.9	328	2.27	28.2	10.9
MultiLayer	PirateNet	105	1151	3.55	—	11.0
MultiLayer	KAN (spline)	97.2	—	—	—	94.7
MultiLayer	WavKAN	63.9	191	47.6	—	99.8
MultiLayer	ChebyshevKAN	23.0	22.8	25.7	9.50	50.2
MultiLayer	FourierWavKAN	13.9	—	—	—	42.1

It is worth being precise about the scope. We are not claiming causal weighting or GradNorm are bad methods; both were introduced to fix real pathologies and are reported to do so. We are claiming something narrower and more awkward: they can silently invert on a problem that looks like the kind they were designed for, and the published practice of reporting a final error without the accompanying schedule diagnostics makes that invisible. A causal-training run computes $w_{\min}$ anyway. If it is pinned at zero, the reported loss is being evaluated on a fraction of the domain, and the number underneath it means something different from what it appears to mean.

Why this problem breaks the schedule is not mysterious either. Causal weighting gates later times behind earlier ones, which is only sound if the earlier times can actually be driven to tolerance. Here the residual near $t = 0$ is dominated by a sharp initial pulse ($\sigma_g = 0.1$) whose second derivatives are large by construction, and the hard-constraint ansatz adds its own $O(\varepsilon^{-2})$ term there. The first slab never converges, so the gate never opens.

SOAP provides a useful stress test of that reading: a second-order optimiser is exactly what one might expect to help on a stiff, badly scaled first slab. It is nevertheless the highest-variance arm in the sweep, so its results are not a recommendation. Complementarity between causal weighting and second-order optimisation remains a hypothesis worth testing at a scale this study does not have; reporting either without the frontier diagnostic hides which one is doing the work.

The classical comparison remains the sobering part. Even at 0.302% the best network’s error is $7.5 \times$ that of a 0.0177 s finite-difference solve and roughly five orders of magnitude behind it in wall-clock. That is the expected outcome on a 1D forward problem with a known coefficient field — the setting where classical discretisation is strongest, and not the setting PINNs are argued for. The point of running it is that it is cheap, and it is the only way to notice that a training recipe has gone backwards.

Limitations. One seed per configuration and no error bars, so differences of a few percent between neighbouring configurations should not be read as real; the order-of-magnitude gaps are. The problem is 1D and forward only. Training time is reported on a shared GPU whose other tenants we do not control, so the network timings are upper bounds on a quiet machine, and the classical timings are single-core CPU — the comparison is between what each family costs as it is normally run, not a FLOP-matched one. The hyperparameters are those specified in the recipe rather than tuned per architecture; a per-architecture search would likely improve the network numbers, though it would have to close two orders of magnitude to change the conclusion. Finally, the causal-stall diagnosis is

specific to this initial condition: a smoother pulse would produce a smaller residual near $t = 0$ and might let the frontier advance.

References

- Sokratis J. Anagnostopoulos, Juan Diego Toscano, Nikolaos Stergiopoulos, and George Em Karniadakis. Residual-based attention and connection to information bottleneck theory in pinns. *arXiv preprint arXiv:2307.00379*, 2023.
- Zavareh Bozorgasl and Hao Chen. Wav-kan: Wavelet kolmogorov-arnold networks. *arXiv preprint arXiv:2405.12832*, 2024.
- James Bradbury, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary, Dougal Maclaurin, George Necula, Adam Paszke, Jake VanderPlas, Skye Wanderman-Milne, and Qiao Zhang. JAX: composable transformations of Python+NumPy programs, 2018. URL <http://github.com/jax-ml/jax>.
- Patrick Kidger and Cristian Garcia. Equinox: neural networks in jax via callable pytrees and filtered transformations, 2021. Differentiable Programming workshop at NeurIPS 2021.
- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *International Conference on Learning Representations*, 2015.
- Dong C. Liu and Jorge Nocedal. On the limited memory bfgs method for large scale optimization. *Mathematical Programming*, 45(1):503–528, 1989.
- Ziming Liu, Yixuan Wang, Sachin Vaidya, Fabian Ruehle, James Halverson, Marin Soljačić, Thomas Y. Hou, and Max Tegmark. Kan: Kolmogorov-arnold networks. *arXiv preprint arXiv:2404.19756*, 2024.
- Gerrit Mur. Absorbing boundary conditions for the finite-difference approximation of the time-domain electromagnetic-field equations. *IEEE Transactions on Electromagnetic Compatibility*, EMC-23(4):377–382, 1981.
- Maziar Raissi, Paris Perdikaris, and George E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019.
- Khemraj Shukla, Juan Diego Toscano, Zhicheng Wang, Zongren Zou, and George Em Karniadakis. A comprehensive and fair comparison between mlp and kan representations for differential equations and operator networks. *arXiv preprint arXiv:2406.02917*, 2024.
- Matthew Tancik, Pratul P. Srinivasan, Ben Mildenhall, Sara Fridovich-Keil, Nithin Raghavan, Utkarsh Singhal, Ravi Ramamoorthi, Jonathan T. Barron, and Ren Ng. Fourier features let networks learn high frequency functions in low dimensional domains. *Advances in Neural Information Processing Systems*, 33, 2020.
- Nikhil Vyas, Depen Morwani, Rosie Zhao, Itai Shapira, David Brandfonbrener, Lucas Janson, and Sham Kakade. Soap: Improving and stabilizing shampoo using adam. *arXiv preprint arXiv:2409.11321*, 2024.
- Sifan Wang, Yujun Teng, and Paris Perdikaris. Understanding and mitigating gradient flow pathologies in physics-informed neural networks. *SIAM Journal on Scientific Computing*, 43(5):A3055–A3081, 2021a.
- Sifan Wang, Hanwen Wang, and Paris Perdikaris. On the eigenvector bias of fourier feature networks: From regression to solving multi-scale pdes with physics-informed neural networks. *Computer Methods in Applied Mechanics and Engineering*, 384:113938, 2021b.
- Sifan Wang, Shyam Sankaran, and Paris Perdikaris. Respecting causality is all you need for training physics-informed neural networks. *arXiv preprint arXiv:2203.07404*, 2022.
- Sifan Wang, Shyam Sankaran, Hanwen Wang, and Paris Perdikaris. An expert’s guide to training physics-informed neural networks. *arXiv preprint arXiv:2308.08468*, 2023.

296 Sifan Wang, Bowen Li, Yuhan Chen, and Paris Perdikaris. Piratenets: Physics-informed deep learning
 297 with residual adaptive networks. *Journal of Machine Learning Research*, 25, 2024.

298 Sifan Wang, Ananyae Kumar Bhartari, Bowen Li, and Paris Perdikaris. Gradient alignment in
 299 physics-informed neural networks: A second order optimization perspective. *arXiv preprint*
 300 *arXiv:2502.00604*, 2025.

301 A Reproducing every number

302 No number in this paper is typed by hand. Two scripts read the result files and emit the tables, the
 303 figures, and a file of \newcommand macros that the abstract and body quote, so the prose cannot drift
 304 away from the tables.

305 A.1 Classical baselines

306 The three solvers are pure NumPy and run on any machine in about six minutes, most of it in the two
 307 spectral reference solves per material:

```
308 python wave/numerical/run_classical_benchmark.py \
309     --out results/classical/classical_benchmark.json
```

310 The material profiles used by the solvers are transcribed into NumPy so that dataset generation does
 311 not require the torch stack; `material_profiles.verify_against_materials()` asserts they
 312 agree with the project’s originals to 10^{-12} and should be run before trusting any number produced
 313 here.

314 **Ground truth.** Chebyshev collocation at $n = 512$, with $n = 384$ run alongside as a self-
 315 convergence check. The two differ by $1.3e-12\%$ on *Homogeneous*, $1.9e-05\%$ on *TwoLayer* and
 316 $1.4e-04\%$ on *MultiLayer* — in every case far below the 0.039% floor the second-order schemes
 317 plateau at, which is what makes the floor measurable rather than an artefact of the reference.

318 **Why not a fine finite-difference reference.** The obvious choice would have been `fd2` at a very
 319 high resolution, and it is what the project used previously. It does not work: `fd2`’s error stops falling
 320 at 0.039%, so scoring against a fine `fd2` solve measures agreement with that floor rather than with
 321 the solution. The floor is identical to three significant figures across three materials whose interior
 322 physics differs completely, which is what identifies it as the boundary condition.

323 A.2 Coordinate networks

324 Runs are enumerated by an integer id; `--list-runs` prints the full mapping. A single run is

```
325 python wave/run_experiment.py --run-id 12 --output-dir results/
```

326 and the ablation arms add `--no-gradnorm` (pins $w_{\text{pde}} = w_{\text{bc}} = w_{\text{ic}} = 1$), `--causal-chunks 1`
 327 (one time slab, so every time is weighted equally), and `--run-tag` so an ablation does not overwrite
 328 its baseline.

329 **Hyperparameters.** Identical across every architecture and every arm, because the point of the
 330 comparison is the training recipe rather than a tuned configuration: 100×100 collocation grid,
 331 15 000 Adam steps at 10^{-3} , up to 9 000 continuation steps while $w_{\text{min}} < 0.9$, then 2 000 L-BFGS
 332 steps. Causal weighting uses $M = 16$ slabs with the tolerance annealed from 0.1 to 1.0; GradNorm
 333 refreshes every 100 steps with smoothing $\alpha = 0.9$. Model selection is on a stationary validation
 334 metric (unweighted per-slab residual plus boundary and initial-condition terms), not the weighted
 335 training loss, which moves as the weights move.

336 **Architecture sizes.** One representative configuration per architecture, chosen so that the baseline,
 337 the SOAP/RBA arms and the ablation arms all use the same network: 3 hidden layers of width 64 for
 338 VanillaPINN, FourierFeature ($n_F = 128$, $\sigma = 3$) and PirateNet; 2 hidden layers of width 16 for the
 339 Kolmogorov–Arnold family (KAN $G = 5$, $k = 3$; WavKAN Mexican hat; ChebyshevKAN degree 5;
 340 FourierWavKAN $n_F = 32$).

341 **Reporting.** Each run writes `l2_errors.json` (errors after the Adam and L-BFGS phases, plus the
342 arm’s configuration) and `causal_convergence.json` (the per-slab weight vector sampled through
343 training, which is what Figure 2 plots). Both are small; the network checkpoints are not needed to
344 regenerate any number.

345 A.3 Rebuilding the paper

```
346 python wave/server/fetch_results.py          # fetch GPU results
347 python paper/neurips2026_workshop/make_tables.py
348 python paper/neurips2026_workshop/make_figures.py
349 pdflatex main && bibtex main && pdflatex main && pdflatex main
```

350 B Implementation notes that changed a result

351 Three implementation choices were arrived at by fixing a failure, and all would silently corrupt the
352 comparison if made differently.

353 **The spectral solver needs the first-order form.** Collocating $u_{tt} = \rho^{-1} \partial_x (E \partial_x u)$ directly is the
354 obvious route and is unstable at the absorbing boundary: advancing u at the end nodes from u^{n-1}
355 decouples them from the interior and the corner modes grow without bound within a few hundred
356 steps. Splitting into velocity–stress makes the radiation condition an exact statement about the
357 incoming Riemann invariant, which can be imposed as a projection at every Runge–Kutta stage. The
358 difference is not accuracy but existence: the second-order form produced NaN, not a worse number.

359 **Resampling can hide spectral convergence.** Solvers choose their own time step from stability
360 and their own spatial grid, so results are resampled onto a shared 201×101 evaluation grid. Doing
361 that with linear interpolation is second-order accurate and puts a floor under the measured error
362 at roughly 10^{-3} — which made the spectral solver look *algebraically* convergent and worse than
363 finite differences. The fix is in two parts: round the step count up until the output times are hit
364 exactly, removing time interpolation from the error budget entirely, and interpolate off the Chebyshev
365 nodes barycentrically rather than linearly. After both, the spectral solver converges as it should. Any
366 comparison of a spectral method against a low-order one on a shared output grid is exposed to this.

367 **Continuation must carry the final iterate.** We run continuation in 3 000-step blocks so training
368 can stop once $w_{\min}$ reaches its target. Reloading the best validation checkpoint between blocks is not
369 equivalent to continuing: frontier progress need not improve that scalar metric, so the next block can
370 replay almost the same trajectory from the same parameters. The runner now passes each block’s
371 final model directly to the next block and restores the best-loss checkpoint only once, before L-BFGS.
372 Result files record this handoff mode, and every causal result reported in the paper was regenerated
373 after the fix.

374 **A note on the KAN implementation.** The spline KAN used here is a direct Equinox implemen-
375 tation of a B-spline edge layer with a SiLU base branch (grid size G , spline order k), rather than a
376 wrapper around an external package. It matches the pykan formulation but does not reproduce that
377 package’s adaptive grid refinement, which would change the parameter shapes mid-run — neither the
378 causal schedule nor L-BFGS tolerates that. The basis is verified to form a partition of unity and to
379 admit the second derivatives the residual needs.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and Section 1 state the two measured claims – that switching off causal weighting and GradNorm improves relative L^2 by 1.8–220× on this problem, and that the second-order classical schemes plateau at 0.039% because of their absorbing boundary – and both are supported by Tables 4 and 1 respectively. The abstract also states explicitly that the ablation does not make PINNs competitive with the classical baseline.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 6 ends with a Limitations paragraph covering the single seed per configuration, the 1D forward-only setting, the untuned per-architecture hyperparameters, the shared GPU the network timings were measured on, and the dependence of the causal-stall diagnosis on this particular initial condition.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper contains no theoretical results; every claim is empirical.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Appendix A gives the exact commands, the run-id scheme, and the complete hyperparameter set, which is identical across every architecture and arm. Appendix B documents the three implementation choices (the velocity–stress formulation, the resampling scheme, and continuation-state handoff) that materially change the reported numbers.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All code is in the accompanying repository, including the three classical solvers, the training runner with the ablation flags, and the scripts that regenerate every table, figure and inline number in this paper from the result files. There is no external dataset: the problem is a PDE with an analytic initial condition, so all data is generated by the solvers themselves.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 4 gives the loss, both weighting schemes, the two initial-condition treatments and the two-phase optimiser schedule; Appendix A gives the collocation grid, step budgets, learning rate, causal and GradNorm settings, per-architecture sizes, and the model-selection criterion. Hyperparameters were not tuned per architecture, which is stated as a limitation.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: Each configuration is a single run at a fixed seed, so no error bars are reported; this is stated in the Limitations paragraph. The claim the paper rests on is an order-of-magnitude effect ($1.8\text{--}220\times$) that is consistent in sign across every architecture and material measured, and the accompanying mechanism (Figure 2) is a deterministic property of the schedule rather than a stochastic outcome. Differences of a few percent between neighbouring configurations are explicitly described as not meaningful.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Section 5 and Appendix A report both sides: the classical solvers are single-core CPU with per-run seconds in Table 1, and the network runs are on NVIDIA H100 NVL GPUs with the measured wall-clock spread quoted in the text and plotted in Figure 1. The Limitations paragraph notes that the GPU was shared with other tenants, so network timings are upper bounds.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: The work is a numerical study of PDE solvers using synthetic data. It involves no human subjects, no personal data, and no released model or dataset artefact.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: The paper is a methodological comparison of numerical and neural-network solvers for a classical PDE. It introduces no deployable system and has no direct societal application; the one practical consequence – that a training diagnostic should be reported alongside final errors – is discussed in Section 6.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases no model, dataset, or scraped content that carries a risk of misuse.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.

- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: The methods reimplemented here (causal weighting, GradNorm, residual-based attention, SOAP, and each architecture) are cited to their original publications in Section 4. The software used – JAX and Equinox – is cited and used under its own open-source licence. No third-party dataset is used.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[N/A\]](#)

Justification: The paper introduces no new dataset or pretrained model. The code that produces the results is documented in Appendix A but is a means of reproducing the study rather than a released asset.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: The paper involves no crowdsourcing and no human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper involves no human subjects, so no IRB review was required.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: Large language models were not part of the research method. The contributions – the classical solver implementations, the training ablation, and the analysis – do not involve an LLM as a component.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.